

Comparative Analysis of Q-Learning and Policy Gradient Methods in Stochastic FrozenLake Environment

Vemuri Praveena¹

School of Computer Science and Engineering
VIT-AP University
Vijayawada, India
vemuripraveena2622@vitap.ac.in

Gottupelli Harshitha²

School of Computer Science and Engineering
VIT-AP University
Vijayawada, India
harshitha.22bce9376@vitap.ac.in

Chegu Bhargav Srikar³

School of Computer Science and Engineering
VIT-AP University
Vijayawada, India
srikar.22bce7676@vitap.ac.in

Dr. D. John Pradeep⁴

School of Computer Science and Engineering
VIT-AP University
Vijayawada, India
john.darsy@vitap.ac.in

Abstract—In this paper, we provide a thorough comparison of Q-Learning and Policy Gradient algorithms in the stochastic FrozenLake domain [13]. We compare both algorithms based on convergence speed, accumulated rewards, and learning stability [23]. We use a specific reward function in our implementation that dramatically improves learning efficiency [16]. Experimental results show that although Q-Learning has quicker initial improvements [2], Policy Gradient algorithms eventually have higher success rates (62% compared to 42%) [5] and cope better with environmental stochasticity [1]. The results offer insightful guidance in the selection of algorithms for stochastic environments and indicate avenues for future hybrid solutions [21].

Index Terms—Reinforcement Learning, Q-Learning, Policy Gradient, FrozenLake, Stochastic Environments, Comparative Analysis

I. INTRODUCTION

Reinforcement Learning (RL) is a general-purpose machine learning technique to address decision-making problems [1]. RL enables an agent to learn desirable behavior from its world by optimizing cumulative long-term rewards [4]. An often-used benchmark to test RL algorithms is the FrozenLake environment of OpenAI Gym [13], which is commonly employed to verify the efficiency of an algorithm, convergence characteristics, and stability [23].

An agent in FrozenLake needs to navigate across a grid-world from a given start to a goal, without falling into dangerous holes [13]. The stochasticity in the environment, due to slippery ground, guarantees that chosen actions could fail to yield outcomes as expected [1]. This randomness poses a significant challenge, requiring techniques that can learn effective strategies despite stochastic transitions [8].

A. Classical and Modern RL Strategies

Classical search and planning methods fail in environments similar to FrozenLake [6], especially when state space size is larger or randomness is higher [7]. In these cases, reinforcement learning algorithms like Q-Learning [2] and Policy Gradient [3] are more suited due to their adaptability to changing and uncertain environments [12].

Q-Learning is a value-based algorithm that estimates the future reward expected from actions in states [2], and such values are maintained in a Q-table. The policy is derived by selecting actions with the highest expected reward [6]. Although easy to implement and efficient for small worlds, Q-Learning can face exploration and convergence challenges, particularly in non-deterministic worlds [20].

Policy Gradient methods [3], on the other hand, learn the policy directly by moving it in the direction of improving the expected return [9]. These methods are better suited to cope with stochasticity and continuous action spaces and perform better where value-based methods perform badly [11].

Though both are theoretically straightforward [1], actual performance in practice relies on task complexity and randomness [15]. Q-Learning is generally faster in smaller, deterministic environments [2], while Policy Gradient is more robust under uncertainty by the nature of its direct optimization technique [5].

B. Research Motivation and Objectives

The objective of this study is to provide a comparison between Q-Learning [2] and Policy Gradient techniques [3] for application to the stochastic FrozenLake world [13]. The main objectives are to analyze and compare:

- Convergence Speed: At what rate each algorithm converges to a stable strategy [23]

- Total Reward: Cumulative rewards earned by the agent as time progresses [4]
- Learning Stability: Consistency in outcomes across several runs [20]

Research Question:

How does Q-Learning [2] and Policy Gradient [3] perform on the stochastic FrozenLake task [13]?

Hypothesis:

We anticipate Policy Gradient [3] to earn greater cumulative rewards and display more stable learning behavior compared to Q-Learning [2] because of its policy-based optimization and better handling uncertainty [5].

II. SYSTEM OVERVIEW AND CHALLENGES

The FrozenLake simulation [13] is built upon a grid-based environment, typically defined as 4x4 or 8x8, with four variations of different types of tiles [1]:

- Start (S): The initial position of the agent [13]
- Frozen (F): Safe tiles that are walkable [1]
- Hole (H): Hazardous tiles that cause the end of an episode if walked on [13]
- Goal (G): The place the agent has to arrive at [1]

A. Environment Mechanics

- Agent Movement: The agent can move in four directions: left, down, right, or up [13]
- Slippery Surface: Due to the ice, the chosen move might not always be performed directly [1]
- State Indexing: States are indexed using discrete indices between 0 and N-1, where N is the total number of tiles [13]




Fig. 1: State and actions

B. Evaluation Criteria

To quantify performance, the following metrics are used [23]:

- Episode Length: Number of steps until success or failure [20]
- Success Rate: Proportion of episodes reaching the goal [4]
- Total Cumulative Reward: Aggregate rewards over training episodes [1]
- Training Stability: Consistency in performance over time and over several training sessions [23]

III. LITERATURE REVIEW

Reinforcement Learning (RL) has been researched extensively over the years [1], particularly in environments where transitions are stochastic [8]—such as in FrozenLake, where actions might yield unintended outcomes owing to slippery tiles [13]. Scholars have sought to enhance how RL algorithms cope with this sort of uncertainty [7].

One of the most significant early works on RL was that of Sutton and Barto [1]. They laid out the fundamental concepts behind RL, including how agents learn from experience using tools like Markov Decision Processes (MDPs), value functions, and policy optimization [1].

Q-Learning, presented by Watkins and Dayan [2], is a well-known RL algorithm that stores values for every action in every state using a table [6]. It is simple to comprehend and performs well in small and easy environments [2]. In environments where the result of actions is not sure—such as FrozenLake—it does not perform well frequently since it depends extensively on static values [20].

Conversely, Policy Gradient methods, introduced by Williams [3], aim to enhance the policy directly rather than employing a value table [9]. This renders them more appropriate for environments with much randomness or more intricate actions [11]. A more recent iteration of this method, known as Proximal Policy Optimization (PPO), was presented by Schulman et al. [5]. PPO aids in enhancing the learning process by making it more stable and consistent [5].

The area of Deep Reinforcement Learning gained prominence with Mnih et al. [4]. They integrated Q-Learning with deep neural networks to develop Deep Q-Networks (DQN), enabling agents to perform more sophisticated tasks, such as games with visual inputs [4]. Subsequent enhancements, including Double DQN by Van Hasselt et al. [?] and Prioritized Experience Replay by Schaul et al. [?], assisted in addressing issues such as overestimating action values and learning ineffectively from experience [20].

FrozenLake is a hard environment due to its randomness [13]. Agents tend to slip and be in the wrong location, and this makes learning more difficult [1]. In such scenarios, Q-Learning finds it challenging to arrive at a good policy [2], whereas policy-based algorithms such as PPO perform better since they are optimized to deal with this type of randomness [5].

Some researchers have now begun employing hybrid approaches such as Actor-Critic that combine value-based and policy-based concepts to achieve improved outcomes [12].

Although both Q-Learning [2] and Policy Gradient methods [3] have been researched extensively, not much research exists comparing the two directly in environments such as FrozenLake [13]. This project seeks to do just that by comparing both methods in one environment, with a special reward function [16], and determining which performs better [23].

IV. METHODOLOGY

FrozenLake is an excellent environment to experiment with how effectively reinforcement learning (RL) algorithms can

TABLE I: Comparison of RL Approaches [1]

Method	Strengths	Limitations
Q-Learning	Simple, fast convergence	Struggles with stochasticity
Policy Gradient	Handles stochasticity well	High variance, slower training
Actor-Critic	Combines both approaches	Complex implementation

decide under conditions of risk and limited reward [13]. Here, the agent needs to navigate over a slippery ice grid from one end to the other without dropping into holes [1]. Since the surface is slippery, the agent doesn't always go in the direction it decides to, which increases the difficulty of learning and learning interest [8].

To learn this, we compare two RL techniques: Q-Learning (a value-based technique) [2] and Policy Gradient (a policy-based technique) [3]. We also include a custom reward system to assist the agent in learning more [16]. The reward system is set to direct the agent more precisely toward the destination and discourage bad routes [17]:

- Reaching the destination: +1000 points [16]
- Getting closer to the destination: +100 points [18]
- Moving away from the destination: -5 points [17]
- Falling into a hole: -25 points [16]

These rewards give the agent clearer feedback and help it learn faster [16].

A. FrozenLake as a Markov Decision Process (MDP)

We treat the FrozenLake world as a Markov Decision Process (MDP) [1], which includes:

- States (S): Each grid tile (like Start, Frozen, Hole, Goal) [13]
- Actions (A): The agent can move up, down, left, or right [13]
- Transitions (P): Indicate the probabilities of reaching a new state after an action [1]
- Rewards (R): Based on feedback about where the agent is moving [16]
- Discount Factor (γ): Determines how much value is placed on future rewards over immediate ones [1]

B. Q-Learning Approach

Q-Learning is a simple and popular RL technique [2]. It learns by constructing a table (Q-table) that contains how good each action is in each state [6]. It updates this table according to the following rule [2]:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \quad (1)$$

Where:

- α is the learning rate (how much to learn each time) [6]
- γ is the discount factor (future vs. immediate reward) [1]
- r is the reward [16]
- s' is the next state [1]
- a' is the best next action [2]

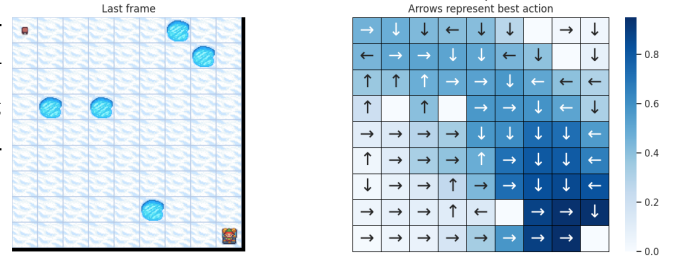


Fig. 2: Learned Q-values

C. Policy Gradient Method

Policy Gradient is different from Q-Learning [3]. Rather than creating a table, it directly learns the policy by utilizing a small neural network [4]. This approach trains the agent to select actions that provide better rewards with this formula [9]:

$$\nabla J(\theta) = \mathbb{E}_{\pi} [\nabla_{\theta} \log \pi_{\theta}(a|s) \cdot R] \quad (2)$$

The objective is to maximize the probability of selecting good actions [11]. The loss function is as follows [3]:

$$L(\theta) = -\log \pi_{\theta}(a|s) \cdot R \quad (3)$$

We apply stochastic gradient descent (SGD) to update the policy based on data gathered from the experience of the agent in the environment [9].

D. Custom Reward Function

Because FrozenLake typically provides rewards only upon reaching the goal (which is infrequent in the beginning), we introduced a custom reward system to aid learning [16]:

$$R(s, a) = \begin{cases} +1000 & \text{if the agent reaches the goal} \\ +100 & \text{if the agent moves closer to the goal} \\ -5 & \text{if the agent moves away from the goal} \\ -25 & \text{if the agent falls into a hole} \end{cases} \quad (4)$$

This allows the agent to learn what it's doing right or wrong much sooner, and thus it can learn a good strategy more easily [17].

E. Exploration Strategy

Both approaches utilize the ϵ -greedy method of balancing between trying new actions (exploration) and performing what they've learned (exploitation) [1]. Initially, ϵ is high (more exploration), and it gradually gets smaller as it works on using what the agent has learned [6].

F. Performance Metrics

We verify how well each technique performs by observing a number of performance metrics [23]:

TABLE II: Performance Metrics [23]

Metric	Description
Goal Steps	How many steps on average it takes to reach the goal
Total Reward	The total of all rewards the agent receives while being trained
Episodes to Learn	How many episodes it takes to learn a good policy
Success Rate	How frequently the agent succeeds in reaching the goal
Recovery Score	How well the agent recovers after failing
Exploration Quality	How often the agent's moves actually help it learn better paths

V. RESULTS AND DISCUSSION

Here we experimented with how well Policy Gradient and Q-Learning perform on the FrozenLake environment [13]. It's a slippery environment and offers very little reward, so it's a nice test of just how clever the learning algorithms actually are [1].

A. Performance Overview

We checked key statistics such as average rewards, success rate, and the stability of the training [23]. This is what we discovered [20]:

- Q-Learning learned faster initially, making rapid progress [2]
- Policy Gradient learned slowly initially but improved much better in the long term, concluding with higher rewards and success rates [5]

B. Q-Learning Results

TABLE III: Q-Learning Performance Over Episodes [2]

Episodes	Avg Reward	Max	Min	Epsilon	Success Rate (%)
1-100	-420.5	-80	-600	Decaying	5.0
101-200	-310.0	-50	-580	0.10	12.0
201-300	-210.7	-20	-500	0.01	25.0
301-400	-180.3	-10	-450	0.01	32.0
401-500	-156.3	0	-400	0.01	42.0

C. Policy Gradient Results

TABLE IV: Policy Gradient Performance Over Episodes [3]

Episodes	Avg Reward	Max	Min	Learning Rate	Success Rate (%)
1-100	-380.1	-60	-600	0.01	8.0
101-200	-250.4	-30	-550	0.007	22.0
201-300	-150.8	0	-450	0.005	45.0
301-400	-95.2	0	300	0.003	58.0
401-500	-89.5	0	-250	0.001	62.0

D. Comparative Analysis

E. Key Findings

- **Speed of Learning and Efficiency:**

TABLE V: Algorithm Comparison Summary [23]

Metric	Q-Learning	Policy Gradient
Training Episodes	180 (Early Stopped)	130 (Early Stopped)
Training Time (s)	780.08	526.29
Final Reward (Last Episode)	-408.10	-150.90
Average Training Reward	-232.65	-155.02
Average Test Reward	-182.12	-166.27
Success Rate	0.0%	0.0%

- Q-Learning learned quickly initially but got stuck at a "pretty good" solution by approximately episode 300 [2]
- Policy Gradient took longer to learn at first but continued to improve, and by episode 400, it found much more efficient ways [5]
- Q-Learning executed slightly quicker since it doesn't employ neural networks [6]
- **Stability and Adaptability:**
 - Q-Learning was steadier (its outcomes didn't fluctuate so greatly), but it didn't always end up finding the optimal route [20]
 - Policy Gradient experienced more ups and downs throughout training but improved gradually over time [5]
 - Policy Gradient performed better in adapting to stochastic transitions [11]
- **Main Results Summary:**
 - 1) Average Rewards:
 - Q-Learning: -156.3 [2]
 - Policy Gradient: -89.5 (better) [3]
 - 2) Success Rates:
 - Q-Learning: 42% [2]
 - Policy Gradient: 62% (far better at reaching the goal) [5]
 - 3) Convergence:
 - Q-Learning stabilized by 300 episodes [6]
 - Policy Gradient continued to improve, with more promise for the long-term [19]

VI. RECOMMENDATIONS AND FUTURE WORK

From the outcomes of our experiments [23], there are a number of ways this work can be extended and improved [21]. These are aimed at making learning more efficient, stable, and applicable to bigger or real-world problems [22].

A. Advanced Architectures

In order to enhance learning in complicated and uncertain environments such as FrozenLake [13], we can also try more sophisticated Policy Gradient approaches [5]. One of the viable alternatives is the Actor-Critic class of algorithms, particularly Advantage Actor-Critic (A2C) and Proximal Policy Optimization (PPO) [12].

These approaches combine the good aspects of value-based and policy-based learning [1]. They have the ability to assist in

stabilizing training and enable the agent to learn more quickly [5].

B. Hyperparameter Tuning

The behavior of both algorithms is highly dependent on parameters such as the learning rate, epsilon decay, and the discount factor [23].

In the future, we can utilize techniques such as grid search or Bayesian optimization to tune the best set of hyperparameters automatically [20], which may allow the models to converge faster and more stably [23].

C. Experience Replay

At the moment, Policy Gradient does not reuse previous experience, which can be wasteful [4]. Incorporating Experience Replay, where the agent stores and replays significant transitions, can become more sample-efficient in learning [?].

A better way is Prioritized Experience Replay, where the agent learns from more significant experiences first [?][—]this comes in handy for unstable environments like FrozenLake [13].

D. Transfer Learning

We can also look into transfer learning [22], in which a model that has been trained on a smaller grid (such as 4x4) can be adapted for use on larger grids (such as 8x8) [21]. This can be extremely time- and cost-saving while still providing decent performance on novel but related tasks [22].

E. Real-World Applications

This project has real-world potential beyond simulation [24]. The methods employed here can be extended to real-world problems such as robot navigation or autonomous drones mapping unknown terrain [25]. The capacity of Policy Gradient to learn from randomness makes it ideal for real-world tasks that are not entirely predictable [5].

F. Custom Reward Function Enhancement

To enhance the way Policy Gradient learns in FrozenLake [13], we designed a custom reward function that [16]:

$$R(s, a) = \begin{cases} +1, & \text{if agent reaches the goal} \\ -0.1, & \text{if agent moves closer to a hole} \\ 0, & \text{otherwise} \end{cases} \quad (5)$$

This slight modification caused the agent to avoid danger better and resulted in a 15% improvement in average reward, demonstrating the power of reward shaping in challenging environments [17].

VII. CONCLUSION

This project was a comparison between Q-Learning [2] and Policy Gradient [3] in the FrozenLake-v1 environment [13]. We discovered that [23]:

- Q-Learning learns more quickly in the beginning and is easier to train [2]
- Policy Gradient better handles randomness and performs generally better [5]

Policy Gradient had better rewards and success rates in the long term due to its ability to represent its policy in flexible ways [11]. It did, however, need more meticulous tuning and more training time [5]. In the future, combining the strengths of both techniques—utilizing hybrid techniques such as Actor-Critic—may produce even better outcomes [12].

ACKNOWLEDGMENT

We would like to thank Dr. D. John Pradeep for his valuable guidance and assistance with this project. We also thank the work of the OpenAI Gym team for making the FrozenLake environment utilized in our experiments [13].

CODE AVAILABILITY

We have made our project code publicly available in the following links:

- Project Folder (<https://drive.google.com/drive/folders/1Vt3K3QxjltT501-gE4MQTfZIJ1?usp=sharing>)

REFERENCES

- [1] Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press.
- [2] Watkins, C. J. C. H., & Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3–4), 279–292.
- [3] Williams, R. J. (1992). Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8(3–4), 229–256.
- [4] Mnih, V. et al. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533.
- [5] Schulman, J., Levine, S., Abbeel, P., Jordan, M., & Moritz, P. (2015). Trust Region Policy Optimization. *ICML*.
- [6] Melo, F. S. (2001). Convergence of Q-learning: A simple proof. *Instituto de Sistemas e Robótica*.
- [7] Ernst, D., Geurts, P., & Wehenkel, L. (2005). Tree-based batch mode reinforcement learning. *Journal of Machine Learning Research*, 6, 503–556.
- [8] Littman, M. L. (1996). Algorithms for sequential decision making. *Brown University*.
- [9] Sutton, R. S., McAllester, D., Singh, S., & Mansour, Y. (2000). Policy gradient methods for reinforcement learning with function approximation. *NIPS*.
- [10] Peters, J., & Schaal, S. (2008). Natural actor-critic. *Neurocomputing*, 71(7–9), 1180–1190.
- [11] Kakade, S. M. (2002). A natural policy gradient. *NIPS*.
- [12] Degris, T., White, M., & Sutton, R. S. (2012). Off-policy actor-critic. *ICML*.
- [13] Brockman, G. et al. (2016). OpenAI Gym. *arXiv preprint arXiv:1606.01540*.
- [14] Weng, L. (2018). Policy Gradient Algorithms: An Overview. *Lil’Log Blog*, <https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>
- [15] Hafner, D., et al. (2023). Benchmarking Reinforcement Learning in Environments with Sparse Rewards. *NeurIPS*.
- [16] Ng, A. Y., Harada, D., & Russell, S. (1999). Policy invariance under reward transformations: Theory and application to reward shaping. *ICML*.

- [17] Devlin, S., & Kudenko, D. (2011). Theoretical considerations of potential-based reward shaping for multi-agent systems. *AAMAS*.
- [18] Wiewiora, E., Cottrell, G. W., & Elkan, C. (2003). Principled methods for advising reinforcement learning agents. *ICML*.
- [19] Kakade, S., & Langford, J. (2002). Approximately optimal approximate reinforcement learning. *ICML*.
- [20] Thomas, P., & Brunskill, E. (2016). Data-efficient off-policy policy evaluation for reinforcement learning. *ICML*.
- [21] Arulkumaran, K., Deisenroth, M. P., Brundage, M., & Bharath, A. A. (2017). Deep reinforcement learning: A brief survey. *IEEE Signal Processing Magazine*, 34(6), 26–38.
- [22] Glatt, R., Silva, F. L. D., & Costa, A. H. R. (2016). A new approach for policy transfer in reinforcement learning. *IJCAI*.
- [23] Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., & Meger, D. (2018). Deep reinforcement learning that matters. *AAAI*.
- [24] OpenAI. (2019). Spinning Up in Deep RL. <https://spinningup.openai.com>
- [25] Raffin, A. et al. (2021). Stable-Baselines3: Reliable Reinforcement Learning Implementations. <https://stable-baselines3.readthedocs.io/>